

Homework 7

Anonymous

2024-02-27

Question 10.1

Using the same crime data set `uscrime.txt` as in Questions 8.2 and 9.1, find the best model you can using

(a) a regression tree model, and

(b) a random forest model.

In R, you can use the `tree` package or the `rpart` package, and the `randomForest` package. For each model, describe one or two qualitative takeaways you get from analyzing the results (i.e., don't just stop when you have a good model, but interpret it too).

Part a

```
library(ggplot2)
library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'


## The following object is masked from 'package:ggplot2':
##
##     margin

library(tree)
crime_data <- read.table(file= "C:\\Users\\sheya\\OneDrive\\Desktop\\uscrime.txt",
                        header = TRUE)
test <- data.frame(M = 14.0, So = 0, Ed = 10.0, Po1 = 12.0, Po2 = 15.5, LF = 0.640,
                  M.F = 94.0, Pop = 150, NW = 1.1, U1 = 0.120, U2 = 3.6, Wealth = 3200,
                  Ineq = 20.1, Prob = 0.04, Time = 39.0)

tree_model <- tree(Crime~., data = crime_data)
summary(tree_model)
```

```
##
## Regression tree:
## tree(formula = Crime ~ ., data = crime_data)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "LF" "NW"
## Number of terminal nodes: 7
## Residual mean deviance: 47390 = 1896000 / 40
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.900 -98.300  -1.545   0.000 110.600  490.100
```

```
#Visulaization
plot(tree_model)
text(tree_model, pretty = 0)
title(main = "Unpruned Regression Tree")
```



```
#Prediction on the unpruned model
pred_un <- predict(tree_model, crime_data[,1:15])


#Calculating R2 for the unpruned model
SSE <- sum((pred_un - crime_data[16])^2)
SST <- sum((crime_data$Crime - mean(crime_data$Crime))^2)
R2 <- 1 - SSE/SST

#Classification trees allow the use of cross-validation to select a good
```

```
#prunning of the tree.
set.seed(18)
tree_cv = cv.tree(tree_model)
summary(tree_cv)
```

```
##      Length Class  Mode
## size    7      -none- numeric
## dev     7      -none- numeric
## k       7      -none- numeric
## method 1      -none- character
```

```
plot(tree_cv$size, tree_cv$dev, type = "b",
      xlab = "Tree Size", ylab = "CV-DEV")
```



```
#The plot shows deviation which is a measurement of the errors from 1 to 7.
#A tree size of 5 and 6 with close proximity of their deviation values show little error.
#Therefore, the unpruned model suggests overfitting.
#Now lets try to prune using tree sizes of 5 and 6 and see the R2.
```

```
#Pruning the tree using a tree size of 6
tree_prune_6 <- prune.tree(tree_model, best = 6)
summary(tree_prune_6)
```

```
##
```

```
## Regression tree:
## snip.tree(tree = tree_model, nodes = 4L)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 6
## Residual mean deviance: 49100 = 2013000 / 41
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.900 -99.520  -1.545    0.000 122.800  490.100
```


```
#Prediction of tree
pred_tree_6 <- predict(tree_prune_6, crime_data[,1:15])

#Calculating R2 for the unpruned model
SSE_cv <- sum((pred_tree_6 - crime_data[16])^2)
R2_6 <- 1 - SSE_cv / SST

#Cross-validation on tree
set.seed(18)
tree_cv_6 = cv.tree(tree_prune_6)
summary(tree_cv_6)
```

```
##      Length Class  Mode
## size    6      -none- numeric
## dev     6      -none- numeric
## k       6      -none- numeric
## method  1      -none- character
```

```
plot(tree_cv_6$size, tree_cv_6$dev, type = 'b',
      xlab = "Tree Size", ylab = "CV-DEV")
```



```
#For each tree will reflect its deviation
```

```
tree_cv_6$size
```

```
## [1] 6 5 4 3 2 1
```

```
tree_cv_6$dev
```

```
## [1] 6912822 7026839 7161659 7674863 6729955 7945960
```

```
#Now lets try for a tree size of 5
```

```
tree_prune_5 <- prune.tree(tree_model, best = 5)
summary(tree_prune_5)
```

```
##
## Regression tree:
## snip.tree(tree = tree_model, nodes = c(4L, 6L))
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 5
## Residual mean deviance: 54210 = 2277000 / 42
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.9 -107.5   15.5     0.0  122.8   490.1
```


```
#Prediction on tree size 5
pred_tree_5 <- predict(tree_prune_5, crime_data[,1:15])
```

```
#Calculating R2 for the unpruned model
SSE_cv <- sum((pred_tree_5 - crime_data[16])^2)
R2_5 <- 1 - SSE_cv / SST
```

```
#Cross-validation on tree
set.seed(18)
tree_cv_5 = cv.tree(tree_prune_5)
summary(tree_cv_5)
```

```
##      Length Class  Mode
## size    5      -none- numeric
## dev     5      -none- numeric
## k       5      -none- numeric
## method 1      -none- character
```

```
plot(tree_cv_5$size, tree_cv_5$dev, type = "b",
      xlab = "Tree Size", ylab = "CV-DEV")
```



```
#For each tree will reflect its deviation
```

```
tree_cv_5$size
```

```
## [1] 5 4 3 2 1
```

```
tree_cv_5$dev
```

```
## [1] 6912822 7161659 7674863 6729955 7945960
```

I used a decision tree to evaluate the small data set. The decision tree was fitted to the full data set called un-pruned data set. The R^2 for this model is 72%, close to what I found with a linear regression model in the previous assignment. I ran cross-validation on the un-pruned data set to spot overfitting. The CV process revealed overfitting which was expected in a small data set. Therefore, pruning was performed to create an optimal decision tree and the same process followed but, this time the R^2 did not decrease with a smaller tree size and overfitting was still pronounced.

Part b

```
#Random Forest model
```

```
rf <- randomForest(Crime~., data = crime_data)
summary(rf)
```

```
##               Length Class  Mode
## call           3      -none- call
## type           1      -none- character
## predicted      47      -none- numeric
## mse            500     -none- numeric
## rsq            500     -none- numeric
## oob.times      47      -none- numeric
## importance     15      -none- numeric
## importanceSD    0      -none- NULL
## localImportance 0      -none- NULL
## proximity      0      -none- NULL
## ntree          1      -none- numeric
## mtry           1      -none- numeric
## forest         11      -none- list
## coefs           0      -none- NULL
## y              47      -none- numeric
## test           0      -none- NULL
## inbag          0      -none- NULL
## terms          3      terms  call
```

```
print(rf)
```

```
##
## Call:
## randomForest(formula = Crime ~ ., data = crime_data)
##               Type of random forest: regression
##               Number of trees: 500
## No. of variables tried at each split: 5
##
##               Mean of squared residuals: 84806.7
##               % Var explained: 42.07
```

```
#Prediction on RF model
```

```
pred_rf <- predict(rf)
```

```
#Calculating R2 for the RF model
```

```
SSE_rf <- sum((pred_rf - crime_data[16])^2)
```

```
SST <- sum((crime_data$Crime - mean(crime_data$Crime))^2)
```

```
R2_rf <- 1 - SSE_rf/SST
```

```
R2_rf
```

```
## [1] 0.42073
```

```
#CV on RF model and the R2
```

```
rf_cv <- rfcv(trainx = crime_data[,1:15], trainy = crime_data$Crime, cv.fold = 10)
```

```
print(rf_cv)
```

```
## $n.var
```

```
## [1] 15 8 4 1
```

```
##
```

```
## $error.cv
```

```
## 15 8 4 1
```

```
## 79809.64 82569.71 83928.17 125964.94
```

```
##
```

```
## $predicted
```

```
## $predicted$'15'
```

```
## [1] 761.7234 1104.4435 687.6527 1361.4968 974.4451 1178.1952 933.9357
```

```
## [8] 1032.6052 842.8781 764.2576 1225.1793 842.1519 708.5982 720.1272
```

```
## [15] 700.5806 893.1612 652.6221 1013.3265 1101.2674 1170.0086 838.4113
```

```
## [22] 717.3869 935.9250 907.6164 681.8581 1264.6006 852.6843 977.2391
```

```
## [29] 1246.1808 809.2042 750.3706 1096.1583 747.8535 937.9947 1097.5014
```

```
## [36] 1060.1829 825.1743 648.5620 787.4838 992.1146 786.3486 587.9489
```

```
## [43] 892.0229 1045.9562 656.7548 1005.5758 1078.4402
```

```
##
```

```
## $predicted$'8'
```

```
## [1] 749.2641 1053.0746 678.1312 1372.5514 969.3615 1275.0023 938.9363
```

```
## [8] 1077.6097 816.8258 723.1820 1253.6596 828.7718 661.9169 678.3712
```

```
## [15] 667.1509 954.2401 596.8021 1065.1728 1206.6320 1247.0461 862.6599
```

```
## [22] 693.2627 926.1622 960.8124 676.5038 1318.9806 838.9362 919.5391
```

```
## [29] 1209.8345 774.7754 710.7930 1139.2371 772.6364 860.5083 1124.6261
```

```
## [36] 1015.2740 793.2516 618.6377 747.5356 941.5504 735.5742 562.5441
```

```
## [43] 910.0046 1049.2343 636.5604 1033.2314 1015.6866
```

```
##
```

```
## $predicted$'4'
```

```
## [1] 716.8167 1057.7331 602.3791 1436.5503 1002.1193 1228.1315 1084.2895
```

```
## [8] 1198.2753 781.3632 707.2577 1122.6516 817.4284 625.4396 680.3395
```

```
## [15] 643.7114 951.9747 629.2351 1254.1266 1308.6034 1326.2521 773.8185
```

```
## [22] 685.5508 965.9302 864.2897 673.2007 1327.2754 817.8120 920.7302
```

```
## [29] 1480.1317 729.3598 672.0886 1068.0030 774.8349 928.7939 1189.8076
```

```
## [36] 1193.2473 775.9754 516.6584 742.3951 995.0158 712.3654 493.3957
```

```
## [43] 927.2906 958.9664 560.8161 1023.3296 930.1657
```

```
##
```

```
## $predicted$'1'
```


```
## [1] 717.4464 667.3418 490.6725 1083.4521 1137.9607 1186.2215 1197.0720
```

```
## [8] 1142.6937 809.9206 841.0243 923.0994 812.1443 599.1325 682.0892
```



```
## [15] 1091.6210 1068.2805 892.4708 1351.0906 1420.9565 1315.9822 878.0493
## [22] 843.0556 846.3730 838.0067 722.4237 1635.8663 615.0410 928.2220
## [29] 1635.8663 792.2858 699.2301 1049.7360 722.4237 599.0559 958.7147
## [36] 1104.9348 843.0556 971.4194 728.6905 928.2220 793.6019 559.1215
## [43] 839.6947 1305.0263 517.7340 1104.9348 954.6451
```

```
plot(rf_cv$error)
```



```
#Prediction using RF_CV model
pred_rf_cv <- rf_cv$predicted[1]

#Calculating R2 for RF_CV model
SSE_rf_cv1 <- sum((pred_rf_cv - crime_data[16])^2)
SST <- sum((crime_data$Crime - mean(crime_data$Crime))^2)
R2_rf_cv <- 1 - SSE_rf_cv1/SST
R2_rf_cv
```

```
## [1] 0.4548623
```

In my findings, the random forest model deals with overfitting sufficiently better than the decision tree in Part a. This can be seen by the errors from the CV model which are close to the initial RF model. Also, the R2 value for the CV model is close to the initial model. Thus, it can be seen that the RF model is good in handling overfitting.

Question 10.2

Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.

Researchers can use predictors that determine the probability of a student getting accepted into a particular university. The relationship between the predictors and the probability of getting accepted is when a logistic regression model is appropriate. The predictors include:

1. GPA
2. ACT score
3. SAT score
4. Number of AP classes
5. Extra curricular activities/ community service hours

The different predictors and the response is either 1 (will be accepted) and 0 (will not be accepted). So, any student can go through the model and its output will either be 1 or 0.

Question 10.3

Part I *Using the GermanCredit data set germancredit.txt use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the glm function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use family=binomial(link="logit") in your glm function call.*

```
library(reshape2)
library(dummy)
```

```
## dummy 0.1.3
```

```
## dummyNews()
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(MASS)
library(car)
```

```
## Loading required package: carData
```

```
library(e1071)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
##
## cov, smooth, var
```

```
german_credit <- read.table(file= "C:\\Users\\sheya\\OneDrive\\Desktop\\germancredit.txt",
                             header = TRUE)
summary(german_credit)
```

```
##      A11              X6              A34              A43
## Length:999      Min.   : 4.00      Length:999      Length:999
## Class :character 1st Qu.:12.00      Class :character Class :character
## Mode  :character Median :18.00      Mode  :character Mode  :character
##                      Mean  :20.92
##                      3rd Qu.:24.00
##                      Max.   :72.00
##      X1169          A65              A75              X4
## Min.   : 250      Length:999      Length:999      Min.   :1.000
## 1st Qu.: 1368      Class :character Class :character 1st Qu.:2.000
## Median : 2320      Mode  :character Mode  :character Median :3.000
## Mean   : 3273
## 3rd Qu.: 3972
## Max.   :18424
##      A93              A101            X4.1              A121
## Length:999      Length:999      Min.   :1.000      Length:999
## Class :character Class :character 1st Qu.:2.000      Class :character
## Mode  :character Mode  :character Median :3.000      Mode  :character
##                      Mean   :2.844
##                      3rd Qu.:4.000
##                      Max.   :4.000
##      X67              A143            A152              X2
## Min.   :19.00      Length:999      Length:999      Min.   :1.000
## 1st Qu.:27.00      Class :character Class :character 1st Qu.:1.000
## Median :33.00      Mode  :character Mode  :character Median :1.000
## Mean   :35.51
## 3rd Qu.:42.00
## Max.   :75.00
##      A173              X1              A192              A201
## Length:999      Min.   :1.000      Length:999      Length:999
## Class :character 1st Qu.:1.000      Class :character Class :character
## Mode  :character Median :1.000      Mode  :character Mode  :character
##                      Mean   :1.155
##                      3rd Qu.:1.000
##                      Max.   :2.000
##      X1.1
## Min.   :1.0
## 1st Qu.:1.0
## Median :1.0
## Mean   :1.3
## 3rd Qu.:2.0
## Max.   :2.0
```

```
#Renaming the columns as per the description in the document
newnames <- c('Checking account Status',
```

```


'Duration in month', 'Credit history',
'Purpose', 'Credit amount',
'Savings account/bonds', 'Employment since',
'Installment rate in % disposable income',
'Status& Sex', 'Other debtors / guarantors',
'Residense since', 'Property', 'Age in years',
'Other installment plans', 'Housing',
'Number of existing credits at this bank', 'Job',
'People being liable to provide maintenance for',
'Telephone', 'foreign worker', 'Customer_Status')
colnames(german_credit) <- newnames

#The document shows there are 7 numeric and 13 categorical predictors
#I will group the numerical variables to visualize the data

num_val <- scale(german_credit[,c(2,5,8,11,13,16,18)])
num_val <- melt(num_val)

#Box plot to view the outliers
ggplot(num_val, aes(x = Var2, y = value)) + geom_boxplot() +
  scale_x_discrete(labels = function(x) stringr::str_wrap(x, width = 10))

```



```

#According to the boxplot, "Credit amount" reveal a significant number of outliers
#However, I cannot remove them for this analysis.

```

```

#The next step so to create dummy variables for the categorical variables
categorical <- german_credit[, -c(2,5,8,11,13,16,18,21)]
numerical <- german_credit[, c(2,5,8,11,13,16,18,21)]

cat_dummy <- dummy(categorical)

#Now I must drop one variable from each category to avoid multi-collinearity.

cat_dummy_new <- cat_dummy[, -c(1,5,10,20,25,30,34,37,41,44,47,51,53)]

#Combining the numerical and dummy data together
new_data <- cbind(cat_dummy_new, numerical)
new_data$Customer_Status <- ifelse(new_data$Customer_Status == 1, 1, 0)

#Response to 1 and 0, 1 = Good and 0 = Bad

table(new_data$Customer_Status)

```

```

##
##    0    1
## 300 699

```

```

set.seed(42)
#Create Data Partition
germ_credit_split <- createDataPartition(y = new_data$Customer_Status,
                                          p = 0.7, times = 1, list = FALSE)
head(germ_credit_split,6)

```

```

##      Resample1
## [1,]         1
## [2,]         2
## [3,]         3
## [4,]         4
## [5,]         7
## [6,]         8

```

```

#Training data
train_data <- new_data[germ_credit_split,]

#Testing data
test_data <- new_data[-germ_credit_split,]

#Initial model with all variables
initial_model <- glm(train_data$Customer_Status~.,
                     family = binomial(link = "logit"),
                     data = train_data)

summary(initial_model)

```

```

##
## Call:
## glm(formula = train_data$Customer_Status ~ ., family = binomial(link = "logit"),

```

```

##      data = train_data)
##
## Coefficients:
##
##              Estimate Std. Error z value
## (Intercept)      2.423e-01  1.391e+00   0.174
## Checking.account.Status_A121      2.869e-01  2.634e-01   1.089
## Checking.account.Status_A131      1.129e+00  5.021e-01   2.249
## Checking.account.Status_A141      1.616e+00  2.840e-01   5.691
## Credit.history_A311      -1.171e-01  7.137e-01  -0.164
## Credit.history_A321      6.137e-01  5.423e-01   1.132
## Credit.history_A331      8.711e-01  5.943e-01   1.466
## Credit.history_A341      1.231e+00  5.508e-01   2.236
## Purpose_A411      1.494e+00  4.526e-01   3.301
## Purpose_A4101      1.073e+00  8.634e-01   1.243
## Purpose_A421      8.633e-01  3.235e-01   2.669
## Purpose_A431      8.176e-01  3.017e-01   2.710
## Purpose_A441      1.101e+00  9.799e-01   1.124
## Purpose_A451      1.636e-01  6.599e-01   0.248
## Purpose_A461     -6.330e-01  4.879e-01  -1.297
## Purpose_A481      1.656e+00  1.197e+00   1.383
## Purpose_A491      3.009e-01  3.954e-01   0.761
## Savings.account.bonds_A621      7.327e-01  3.556e-01   2.060
## Savings.account.bonds_A631      2.318e-01  4.498e-01   0.515
## Savings.account.bonds_A641      1.191e+00  6.289e-01   1.894
## Savings.account.bonds_A651      9.366e-01  3.368e-01   2.781
## Employment.since_A721     -3.904e-01  5.289e-01  -0.738
## Employment.since_A731      1.650e-01  5.054e-01   0.326
## Employment.since_A741      6.894e-01  5.439e-01   1.267
## Employment.since_A751      3.031e-01  5.096e-01   0.595
## Status..Sex_A921      3.689e-01  4.449e-01   0.829
## Status..Sex_A931      7.688e-01  4.357e-01   1.765
## Status..Sex_A941      6.067e-01  5.484e-01   1.106
## Other.debtors...guarantors_A1021     -1.142e+00  4.906e-01  -2.329
## Other.debtors...guarantors_A1031      7.616e-01  5.020e-01   1.517
## Property_A1221     -3.140e-01  3.087e-01  -1.017
## Property_A1231     -1.904e-01  2.815e-01  -0.676
## Property_A1241     -1.740e-01  5.246e-01  -0.332
## Other.installment.plans_A1421      1.144e-01  4.937e-01   0.232
## Other.installment.plans_A1431      6.186e-01  3.000e-01   2.062
## Housing_A1521      4.877e-01  2.786e-01   1.750
## Housing_A1531     -1.525e-02  5.937e-01  -0.026
## Job_A1721     -6.614e-01  9.171e-01  -0.721
## Job_A1731     -1.011e+00  8.887e-01  -1.137
## Job_A1741     -8.998e-01  8.998e-01  -1.000
## Telephone_A1921      2.877e-01  2.525e-01   1.139
## foreign.worker_A2021      1.750e+00  7.470e-01   2.343
## 'Duration in month'     -3.402e-02  1.132e-02  -3.007
## 'Credit amount'     -1.288e-04  5.519e-05  -2.333
## 'Installment rate in % disposable income'     -3.578e-01  1.098e-01  -3.260
## 'Residence since'     -3.172e-02  1.035e-01  -0.306
## 'Age in years'      1.906e-02  1.133e-02   1.683
## 'Number of existing credits at this bank'     -1.935e-01  2.233e-01  -0.866
## 'People being liable to provide maintenance for'     -2.890e-01  3.054e-01  -0.946
##
## Pr(>|z|)

```

```

## (Intercept) 0.861669
## Checking.account.Status_A121 0.276038
## Checking.account.Status_A131 0.024518 *
## Checking.account.Status_A141 1.27e-08 ***
## Credit.history_A311 0.869614
## Credit.history_A321 0.257807
## Credit.history_A331 0.142684
## Credit.history_A341 0.025378 *
## Purpose_A411 0.000962 ***
## Purpose_A4101 0.213958
## Purpose_A421 0.007616 **
## Purpose_A431 0.006722 **
## Purpose_A441 0.260977
## Purpose_A451 0.804224
## Purpose_A461 0.194522
## Purpose_A481 0.166560
## Purpose_A491 0.446675
## Savings.account.bonds_A621 0.039367 *
## Savings.account.bonds_A631 0.606357
## Savings.account.bonds_A641 0.058279 .
## Savings.account.bonds_A651 0.005427 **
## Employment.since_A721 0.460405
## Employment.since_A731 0.744072
## Employment.since_A741 0.205019
## Employment.since_A751 0.551987
## Status..Sex_A921 0.407027
## Status..Sex_A931 0.077640 .
## Status..Sex_A941 0.268563
## Other.debtors...guarantors_A1021 0.019880 *
## Other.debtors...guarantors_A1031 0.129192
## Property_A1221 0.309166
## Property_A1231 0.498846
## Property_A1241 0.740054
## Other.installment.plans_A1421 0.816690
## Other.installment.plans_A1431 0.039191 *
## Housing_A1521 0.080074 .
## Housing_A1531 0.979506
## Job_A1721 0.470781
## Job_A1731 0.255447
## Job_A1741 0.317323
## Telephone_A1921 0.254655
## foreign.worker_A2021 0.019133 *
## 'Duration in month' 0.002641 **
## 'Credit amount' 0.019649 *
## 'Installment rate in % disposable income' 0.001116 **
## 'Residence since' 0.759277
## 'Age in years' 0.092459 .
## 'Number of existing credits at this bank' 0.386261
## 'People being liable to provide maintenance for' 0.344055
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##

```

```
## Null deviance: 868.33 on 699 degrees of freedom
## Residual deviance: 621.17 on 651 degrees of freedom
## AIC: 719.17
##
## Number of Fisher Scoring iterations: 5
```

*#From the summary, there are plenty non-significant p-values.
#I will build a new model with only significant variables and a threshold of 0.05.*

```
new_model <- glm(train_data$Customer_Status ~ Checking.account.Status_A13 +
  Checking.account.Status_A14 + Credit.history_A34 +
  Purpose_A41 + Purpose_A42 + Purpose_A43 +
  Savings.account.bonds_A62 + Savings.account.bonds_A65 +
  Other.debtors...guarantors_A102 + foreign.worker_A202 +
  `Duration in month` + `Credit amount` +
  `Installment rate in % disposable income`,
  family = binomial(link = "logit"), data = train_data)

summary(new_model)
```

```
##
## Call:
## glm(formula = train_data$Customer_Status ~ Checking.account.Status_A13 +
##   Checking.account.Status_A14 + Credit.history_A34 + Purpose_A41 +
##   Purpose_A42 + Purpose_A43 + Savings.account.bonds_A62 + Savings.account.bonds_A65 +
##   Other.debtors...guarantors_A102 + foreign.worker_A202 + `Duration in month` +
##   `Credit amount` + `Installment rate in % disposable income`,
##   family = binomial(link = "logit"), data = train_data)
##
## Coefficients:
##
##               Estimate Std. Error z value
## (Intercept)      1.518e+00  3.851e-01   3.941
## Checking.account.Status_A131      8.649e-01  4.424e-01   1.955
## Checking.account.Status_A141      1.504e+00  2.242e-01   6.710
## Credit.history_A341      7.287e-01  2.273e-01   3.206
## Purpose_A411      1.140e+00  3.855e-01   2.958
## Purpose_A421      4.411e-01  2.573e-01   1.715
## Purpose_A431      7.852e-01  2.386e-01   3.291
## Savings.account.bonds_A621      6.421e-01  3.269e-01   1.964
## Savings.account.bonds_A651      9.959e-01  3.024e-01   3.293
## Other.debtors...guarantors_A1021     -1.193e+00  4.670e-01  -2.554
## foreign.worker_A2021      1.375e+00  7.131e-01   1.928
## `Duration in month`     -4.032e-02  1.033e-02  -3.903
## `Credit amount`     -7.889e-05  4.802e-05  -1.643
## `Installment rate in % disposable income` -2.832e-01  9.758e-02  -2.902
##
##               Pr(>|z|)
## (Intercept)      8.10e-05 ***
## Checking.account.Status_A131      0.050544 .
## Checking.account.Status_A141      1.94e-11 ***
## Credit.history_A341      0.001345 **
## Purpose_A411      0.003097 **
## Purpose_A421      0.086384 .
## Purpose_A431      0.000998 ***
## Savings.account.bonds_A621      0.049514 *
```



```
## Savings.account.bonds_A651          0.000990 ***
## Other.debtors...guarantors_A1021     0.010642 *
## foreign.worker_A2021                 0.053836 .
## 'Duration in month'                  9.50e-05 ***
## 'Credit amount'                     0.100368
## 'Installment rate in % disposable income' 0.003706 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 868.33  on 699  degrees of freedom
## Residual deviance: 685.45  on 686  degrees of freedom
## AIC: 713.45
##
## Number of Fisher Scoring iterations: 5
```

```
#This new model is my final model as most of the variables are significant.
#I will need to change the test data with the same variables as the ones in the new model
#for predictions
```

```
test_new<-test_data[,c('Checking.account.Status_A13',
                       'Checking.account.Status_A14','Credit.history_A34',
                       'Purpose_A41','Purpose_A42','Purpose_A43',
                       'Savings.account.bonds_A62','Savings.account.bonds_A65',
                       'Other.debtors...guarantors_A102',
                       'Other.installment.plans_A143',
                       'foreign.worker_A202','Duration in month','Credit amount',
                       'Installment rate in % disposable income','Customer_Status')]
```

```
#Accuracy defines how effective the model is in characterizing bad and good status.
```

```
#Calculate the predicted probabilities for the test data using 0.5 as threshold
predicted <- predict(new_model, test_new, type = "response")
predicted_int0.5 <- as.integer(predicted > 0.5)
```

```
table <- table(predicted_int0.5, test_new$Customer_Status)
confusion <- confusionMatrix(table, positive = '1')
confusion
```

```
## Confusion Matrix and Statistics
##
##
## predicted_int0.5    0    1
##                   0  34  26
##                   1  48 191
##
##               Accuracy : 0.7525
##               95% CI : (0.6996, 0.8004)
##      No Information Rate : 0.7258
##      P-Value [Acc > NIR] : 0.16558
##
##               Kappa : 0.3217
##
```

```
## McNemar's Test P-Value : 0.01464
##
##      Sensitivity : 0.8802
##      Specificity : 0.4146
##      Pos Pred Value : 0.7992
##      Neg Pred Value : 0.5667
##      Prevalence : 0.7258
##      Detection Rate : 0.6388
##      Detection Prevalence : 0.7993
##      Balanced Accuracy : 0.6474
##
##      'Positive' Class : 1
##
```

#Accuracy predicted with 0.5 threshold is 75%

```
#I would like to try 0.6 as a threshold to see the accuracy
predicted_int0.6 <- as.integer(predicted > 0.6)
table2 <- table(predicted_int0.6, test_new$Customer_Status)
confusion2 <- confusionMatrix(table2)
confusion2
```

```
## Confusion Matrix and Statistics
##
##
## predicted_int0.6    0    1
##                   0  46  45
##                   1  36 172
##
##      Accuracy : 0.7291
##      95% CI : (0.6749, 0.7787)
##      No Information Rate : 0.7258
##      P-Value [Acc > NIR] : 0.4780
##
##      Kappa : 0.3419
##
## McNemar's Test P-Value : 0.3741
##
##      Sensitivity : 0.5610
##      Specificity : 0.7926
##      Pos Pred Value : 0.5055
##      Neg Pred Value : 0.8269
##      Prevalence : 0.2742
##      Detection Rate : 0.1538
##      Detection Prevalence : 0.3043
##      Balanced Accuracy : 0.6768
##
##      'Positive' Class : 0
##
```


#Accuracy predicted with 0.6 threshold is 71%

```
#ROC
roc_curve <- roc(test_new$Customer_Status, predicted)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve)
```



```
roc_curve$auc
```

```
## Area under the curve: 0.7538
```

*#Therefore, the accuracy is best between 0.5 and 0.6 thresholds since the ROC curve
#reveals the AUC measuring separability.
#Higher the AUC, the better the model is at predicting 0 classes as 0
#and 1 classes as 1.*

Part II Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between “good” and “bad” answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.

*#This part asks for the number of loans to false positives are reduced.
#Therefore, the sensitivity needs to be high
#This is accomplished by testing different thresholds and seeing where the*

```
#sensitivity is highest
```

```
#Previously I tested a threshold of 0.6
```

```
predicted_int0.6 <- as.integer(predicted > 0.6)
table2 <- table(predicted_int0.6, test_new$Customer_Status)
confusion2 <- confusionMatrix(table2, positive = '1')
confusion2
```

```
## Confusion Matrix and Statistics
##
##
## predicted_int0.6    0    1
##                   0  46  45
##                   1  36 172
##
##               Accuracy : 0.7291
##               95% CI : (0.6749, 0.7787)
##      No Information Rate : 0.7258
##      P-Value [Acc > NIR] : 0.4780
##
##               Kappa : 0.3419
##
##  Mcnemar's Test P-Value : 0.3741
##
##               Sensitivity : 0.7926
##               Specificity : 0.5610
##               Pos Pred Value : 0.8269
##               Neg Pred Value : 0.5055
##               Prevalence : 0.7258
##               Detection Rate : 0.5753
##      Detection Prevalence : 0.6957
##      Balanced Accuracy : 0.6768
##
##      'Positive' Class : 1
##
```

```
#The sensitivity predicted 79% with a threshold of 0.6
```

```
#Threshold of 0.4
```

```
predicted_int0.4 <- as.integer(predicted > 0.4)
table3 <- table(predicted_int0.4, test_new$Customer_Status)
confusion3 <- confusionMatrix(table3, positive = '1')
confusion3
```

```
## Confusion Matrix and Statistics
##
##
## predicted_int0.4    0    1
##                   0  18  16
##                   1  64 201
##
##               Accuracy : 0.7324
##               95% CI : (0.6784, 0.7818)
```

```
##      No Information Rate : 0.7258
##      P-Value [Acc > NIR] : 0.4266
##
##              Kappa : 0.1782
##
##      McNemar's Test P-Value : 1.482e-07
##
##              Sensitivity : 0.9263
##              Specificity : 0.2195
##              Pos Pred Value : 0.7585
##              Neg Pred Value : 0.5294
##              Prevalence : 0.7258
##              Detection Rate : 0.6722
##      Detection Prevalence : 0.8863
##      Balanced Accuracy : 0.5729
##
##      'Positive' Class : 1
##
```

#The sensitivity predicted 92% with a threshold of 0.4

#threshold of 0.2

```
predicted_int0.2 <- as.integer(predicted > 0.2)
table4 <- table(predicted_int0.2, test_new$Customer_Status)
confusion4 <- confusionMatrix(table4, positive = '1')
confusion4
```

```
## Confusion Matrix and Statistics
##
##
## predicted_int0.2    0    1
##                0    7    2
##                1   75  215
##
##              Accuracy : 0.7425
##              95% CI : (0.689, 0.7911)
##      No Information Rate : 0.7258
##      P-Value [Acc > NIR] : 0.2821
##
##              Kappa : 0.1053
##
##      McNemar's Test P-Value : 2.303e-16
##
##              Sensitivity : 0.99078
##              Specificity : 0.08537
##              Pos Pred Value : 0.74138
##              Neg Pred Value : 0.77778
##              Prevalence : 0.72575
##              Detection Rate : 0.71906
##      Detection Prevalence : 0.96990
##      Balanced Accuracy : 0.53807
##
##      'Positive' Class : 1
##
```

```
#The sensitivity predicted 98% with a threshold of 0.2
```

```
#Threshold 0.1
```

```
predicted_int0.1 <- as.integer(predicted > 0.1)
table5 <- table(predicted_int0.1, test_new$Customer_Status)
confusion5 <- confusionMatrix(table5, positive = '1')
confusion5
```

```
## Confusion Matrix and Statistics
##
##
## predicted_int0.1    0    1
##                   0    2    2
##                   1   80  215
##
##               Accuracy : 0.7258
##               95% CI   : (0.6714, 0.7755)
##      No Information Rate : 0.7258
##      P-Value [Acc > NIR] : 0.5297
##
##               Kappa   : 0.0216
##
##  Mcnemar's Test P-Value : <2e-16
##
##               Sensitivity : 0.99078
##               Specificity : 0.02439
##               Pos Pred Value : 0.72881
##               Neg Pred Value : 0.50000
##               Prevalence : 0.72575
##               Detection Rate : 0.71906
##      Detection Prevalence : 0.98662
##      Balanced Accuracy : 0.50759
##
##      'Positive' Class : 1
##
```

```
#The sensitivity predicted 98% with a threshold of 0.1
```

Based on my findings, the threshold increases the sensitivity as the threshold lowers. However, there was no significant change in sensitivity when the threshold was tested below 0.2. In conclusion, the threshold of 0.2 is a good probability based on the model performed.